

In [5]:

```
1 import warnings
2 import pandas as pd
3 import seaborn as sns
4 import matplotlib.pyplot as plt
5 from category_encoders import OneHotEncoder
6 from ipywidgets import Dropdown, FloatSlider, IntSlider, interact
7 from sklearn.impute import SimpleImputer
8 from sklearn.metrics import mean_absolute_error
9 from sklearn.pipeline import make_pipeline
10 warnings.simplefilter(action="ignore", category=FutureWarning)
```

EDA Perfoming

Data Cleaning

```
In [18]: 1 def wrangle(filepath):
2         # Read CSV file
3         df = pd.read_csv(filepath)
4
5         # Subset data: Remove outliers for "a"
6         q1 = df["a"].quantile(0.25)
7         q3 = df["a"].quantile(0.75)
8         iqr = q3 - q1
9         lower_bound = q1 - (1.5 * iqr)
10        upper_bound = q3 + (1.5 * iqr)
11        df = df[(df["a"] > lower_bound) & (df["a"] < upper_bound)]
12
13        # Subset data: Remove outliers for "rms"
14        q1 = df["rms"].quantile(0.25)
15        q3 = df["rms"].quantile(0.75)
16        iqr = q3 - q1
17        lower_bound = q1 - (1.5 * iqr)
18        upper_bound = q3 + (1.5 * iqr)
19        df = df[(df["rms"] > lower_bound) & (df["rms"] < upper_bound)]
20
21
22        # Subset data: Remove outliers for "i"
23        q1 = df["i"].quantile(0.25)
24        q3 = df["i"].quantile(0.75)
25        iqr = q3 - q1
26        lower_bound = q1 - (1.5 * iqr)
27        upper_bound = q3 + (1.5 * iqr)
28        df = df[(df["i"] > lower_bound) & (df["i"] < upper_bound)]
29
30        # Subset data: Remove outliers for "diameter"
31        q1 = df["diameter"].quantile(0.25)
32        q3 = df["diameter"].quantile(0.75)
33        iqr = q3 - q1
34        lower_bound = q1 - (1.5 * iqr)
35        upper_bound = q3 + (1.5 * iqr)
36        df = df[(df["diameter"] > lower_bound) & (df["diameter"] < upper_bound)]
37
38
39        # dropping columns having no correlation with diameter
40        df.drop(columns = ["om", "w"] , inplace = True)
41
42        #dropping low and high cardinality categorical columns
43        df.drop(columns = ["full_name", "classes", "orbit_id", "producer"], inplace = True)
44
45
46        #dropping multicollinearity
47        df.drop(columns=["first_year_obs", "moid_jup"], inplace=True)
48
49        return df
```

```
In [19]: 1 df = wrangle("top_asteroids.csv")
          2 df.shape
```

Out[19]: (113442, 16)

```
In [4]: 1 df.head()
```

Out[4]:

	e	a	i	om	w	ma	n	tp
2	0.160543	2.228812	1.747387	121.579382	252.465454	208.942016	0.296206	2.459110e+06
3	0.167945	2.241299	2.428619	161.636895	172.846491	20.350289	0.293734	2.458531e+06
13	0.278983	2.545744	12.715483	357.019751	349.524955	156.468717	0.242651	2.457956e+06
14	0.137261	2.174837	2.458112	213.756930	174.868642	291.496345	0.307301	2.458823e+06
15	0.107869	2.180337	4.273327	281.903596	90.788273	301.127993	0.306139	2.458793e+06

Features with High and Low Cardinality

```
In [20]: 1 df["classes"].value_counts()
```

Out[20]:

MBA	118581
OMB	6122
IMB	567
APO	429
MCA	342
AMO	224
TJN	130
ATE	87
CEN	11
AST	2
TNO	2

Name: classes, dtype: int64

```
In [22]: 1 df["full_name"].value_counts().head()
```

Out[22]:

266455 (2007 JH42)	1
131803 (2002 AD59)	1
106082 (2000 SP350)	1
16259 Housinger (2000 JR13)	1
397146 (2005 WM162)	1

Name: full_name, dtype: int64

In [23]: 1 df["orbit_id"].value_counts()

Out[23]:

JPL 15	8631
JPL 16	8448
JPL 14	8414
JPL 17	7934
JPL 18	7802
...	
JPL 223	1
JPL 140	1
JPL 712	1
JPL 236	1
JPL 134	1

Name: orbit_id, Length: 222, dtype: int64

In [24]: 1 df["producer"].value_counts()

Out[24]:

Otto Matic	126490
Davide Farnocchia	4
Giorgini	2
Ryan S. Park	1

Name: producer, dtype: int64

In [9]: 1 df.select_dtypes(exclude="number").head()

Out[9]:

	full_name	orbit_id	classes	producer
0	228 Agathe	JPL 35	MBA	Otto Matic
1	290 Bruna	JPL 25	MBA	Otto Matic
2	296 Phaetusa	JPL 28	MBA	Otto Matic
3	315 Constantia	JPL 35	MBA	Otto Matic
4	330 Adalberta (A910 CB)	JPL 34	MBA	Otto Matic

In [30]: 1 df.info()

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 126497 entries, 0 to 126496
Data columns (total 18 columns):
#   Column                Non-Null Count  Dtype  
---  -
0   e                      126497 non-null float64
1   a                      126497 non-null float64
2   i                      126497 non-null float64
3   ma                     126497 non-null float64
4   n                      126497 non-null float64
5   tp                     126497 non-null float64
6   moid                   126497 non-null float64
7   moid_jup               126497 non-null float64
8   data_arc               126497 non-null float64
9   n_obs_used             126497 non-null int64  
10  rms                     126497 non-null float64
11  diameter               126497 non-null float64
12  albedo                 126497 non-null float64
13  diameter_sigma         126497 non-null float64
14  first_year_obs         126497 non-null int64  
15  first_month_obs        126497 non-null int64  
16  last_obs_year          126497 non-null int64  
17  last_obs_month         126497 non-null int64  
dtypes: float64(13), int64(5)
memory usage: 17.4 MB

```

In [7]: 1 df.describe()

Out[7]:

	e	a	i	om	w	m
count	126497.000000	126497.000000	126497.000000	126497.000000	126497.000000	126497.000000
mean	0.146644	2.756965	10.203665	169.819406	181.823887	182.53216
std	0.076841	0.453027	6.689924	102.749965	103.538522	103.41604
min	0.000488	0.626226	0.021855	0.000929	0.004466	0.00051
25%	0.091182	2.510297	5.051481	82.100534	91.822257	93.74634
50%	0.140047	2.729370	9.244113	160.539684	183.660501	185.54257
75%	0.192297	3.074005	13.538838	256.258893	271.540490	270.95750
max	0.968381	69.576833	158.535394	359.990858	359.995174	359.99922

In [138]: 1 df.corr()

Out[138]:

	e	a	i	ma	n	tp	moid	data
e	1.000000	-0.152621	0.093734	-0.007079	0.145126	0.009181	-0.645217	-0.011
a	-0.152621	1.000000	0.242624	0.080987	-0.988474	0.141626	0.847223	-0.22
i	0.093734	0.242624	1.000000	0.009375	-0.232147	-0.063812	0.172244	-0.22
ma	-0.007079	0.080987	0.009375	1.000000	-0.070111	0.357760	0.062824	-0.01
n	0.145126	-0.988474	-0.232147	-0.070111	1.000000	-0.143300	-0.832809	0.22
tp	0.009181	0.141626	-0.063812	0.357760	-0.143300	1.000000	0.099002	0.15
moid	-0.645217	0.847223	0.172244	0.062824	-0.832809	0.099002	1.000000	-0.16
data_arc	-0.013686	-0.221385	-0.222807	-0.018713	0.221875	0.157837	-0.168356	1.00
n_obs_used	-0.073396	-0.321318	-0.224627	-0.044744	0.318419	0.066747	-0.214607	0.67
rms	-0.086941	0.209375	-0.005380	0.025502	-0.210185	0.054743	0.204195	-0.33
diameter	-0.125871	0.505332	0.148615	0.013704	-0.493851	0.079103	0.457017	0.27
albedo	-0.030027	-0.432693	-0.055693	-0.042709	0.437851	0.025229	-0.317046	0.35
diameter_sigma	-0.063755	0.231942	0.038652	-0.017350	-0.230362	-0.049323	0.212657	-0.13
first_month_obs	0.029050	0.007217	-0.028954	0.005383	-0.008671	0.061897	-0.009492	0.08
last_obs_year	-0.068083	0.050000	-0.132369	-0.007361	-0.065980	0.566791	0.070338	0.35
last_obs_month	-0.038546	-0.014067	-0.008514	0.082184	0.016243	0.099939	0.010149	0.12


Corelation between independend features

In [139]: 1 df.drop(columns=["diameter"]).corr()

Out[139]:

	e	a	i	ma	n	tp	moid	data
e	1.000000	-0.152621	0.093734	-0.007079	0.145126	0.009181	-0.645217	-0.011
a	-0.152621	1.000000	0.242624	0.080987	-0.988474	0.141626	0.847223	-0.22
i	0.093734	0.242624	1.000000	0.009375	-0.232147	-0.063812	0.172244	-0.22
ma	-0.007079	0.080987	0.009375	1.000000	-0.070111	0.357760	0.062824	-0.01
n	0.145126	-0.988474	-0.232147	-0.070111	1.000000	-0.143300	-0.832809	0.22
tp	0.009181	0.141626	-0.063812	0.357760	-0.143300	1.000000	0.099002	0.15
moid	-0.645217	0.847223	0.172244	0.062824	-0.832809	0.099002	1.000000	-0.16
data_arc	-0.013686	-0.221385	-0.222807	-0.018713	0.221875	0.157837	-0.168356	1.00
n_obs_used	-0.073396	-0.321318	-0.224627	-0.044744	0.318419	0.066747	-0.214607	0.67
rms	-0.086941	0.209375	-0.005380	0.025502	-0.210185	0.054743	0.204195	-0.33
albedo	-0.030027	-0.432693	-0.055693	-0.042709	0.437851	0.025229	-0.317046	0.35
diameter_sigma	-0.063755	0.231942	0.038652	-0.017350	-0.230362	-0.049323	0.212657	-0.13
first_month_obs	0.029050	0.007217	-0.028954	0.005383	-0.008671	0.061897	-0.009492	0.08
last_obs_year	-0.068083	0.050000	-0.132369	-0.007361	-0.065980	0.566791	0.070338	0.35
last_obs_month	-0.038546	-0.014067	-0.008514	0.082184	0.016243	0.099939	0.010149	0.12

```
In [6]: 1 #correlation matrix
2 corrmat = df.drop(columns=["diameter"]).corr()
3 f, ax = plt.subplots(figsize=(12, 9))
4 sns.heatmap(corrmat, vmax=.8, square=True)
5 plt.title("correlation matrix of independent variables")
6 plt.savefig("correlation_matrix.png")
```



```
In [71]: 1 # print("correlation between diameter and first year observed",df["diamete
2 print("correlation between diameter and data_arc",df["diameter"].corr(df["
3 # print("correlation between diameter and n",df["diameter"].corr(df["n"]))
4 print("correlation between diameter and a",df["diameter"].corr(df["a"]))
5 print("correlation between diameter and albedo",df["diameter"].corr(df["al
6 print("correlation between diameter and moid_jup",df["diameter"].corr(df["
7 print("correlation between diameter and e",df["diameter"].corr(df["e"]))
```


```
correlation between diameter and data_arc 0.30657721130460047
correlation between diameter and a 0.4960066880908689
correlation between diameter and albedo -0.26085838909684234
correlation between diameter and moid_jup 0.45273686921269474
correlation between diameter and e -0.1224700301510987
```



```

In [7]: 1 #correlation matrix
2 corrmat = df.corr()
3 f, ax = plt.subplots(figsize=(12, 9))
4 sns.heatmap(corrmat, vmax=.8, square=True)
5 plt.title("correlation matrix of independent variables with Dimeter")
6 plt.savefig("correlation_matrix_with_target.png")


```



```

In [15]: 1 # Create a figure with three subplots
2 fig, axes = plt.subplots(nrows=1, ncols=3, figsize=(18, 6))
3
4 # Plot 1: 'e' vs 'a'
5 axes[0].scatter(df['e'], df['a'])
6 axes[0].set_xlabel('e')
7 axes[0].set_ylabel('a')
8 axes[0].set_title('Scatter Plot: e vs a')
9
10 # Plot 2: 'a' vs 'diameter'
11 axes[1].scatter(df['a'], df['diameter'])
12 axes[1].set_xlabel('a')
13 axes[1].set_ylabel('diameter')
14 axes[1].set_title('Scatter Plot: a vs diameter')
15
16 # Plot 3: 'albedo' vs 'diameter'
17 axes[2].scatter(df['albedo'], df['diameter'])
18 axes[2].set_xlabel('albedo')
19 axes[2].set_ylabel('diameter')
20 axes[2].set_title('Scatter Plot: albedo vs diameter')
21
22 # Adjust spacing between subplots
23 plt.tight_layout()
24
25 # Save the plot as a PNG file
26 plt.savefig('side_by_side_plots.png')
27
28 # Display the plot
29 plt.show()

```




Outlier Analysis

Diameter Before removing outlier

```
In [150]: 1 sns.histplot(data=df, x="diameter", bins=30)
```


```
Out[150]: <AxesSubplot:xlabel='diameter', ylabel='Count'>
```



After Removing Outlier it is almost normaly distributed


```
In [155]: 1 sns.histplot(data=df, x="diameter", bins=30)
```

```
Out[155]: <AxesSubplot:xlabel='diameter', ylabel='Count'>
```




```
In [149]: 1 plt.figure(figsize=(8,4))  
          2 sns.boxplot(data=df[["diameter"]])
```

Out[149]: <AxesSubplot:>




```
In [156]: 1 plt.figure(figsize=(8,4))  
          2 sns.boxplot(data=df[["diameter"]])
```

Out[156]: <AxesSubplot:>




```
In [75]: 1 plt.figure(figsize=(8,4))
          2 sns.boxplot(data=df[["rms"]])
```

Out[75]: <AxesSubplot:>




```
In [79]: 1 plt.figure(figsize=(8,4))
          2 sns.boxplot(data=df[["rms"]])
```

Out[79]: <AxesSubplot:>




```
In [39]: 1 plt.figure(figsize=(12,8))  
2 sns.boxplot(data=df[['e', 'a', 'i']])
```

Out[39]: <AxesSubplot:>



```
In [94]: 1 plt.figure(figsize=(12,8))  
2 sns.boxplot(data=df[['e', 'a', 'i', "moid"]])
```

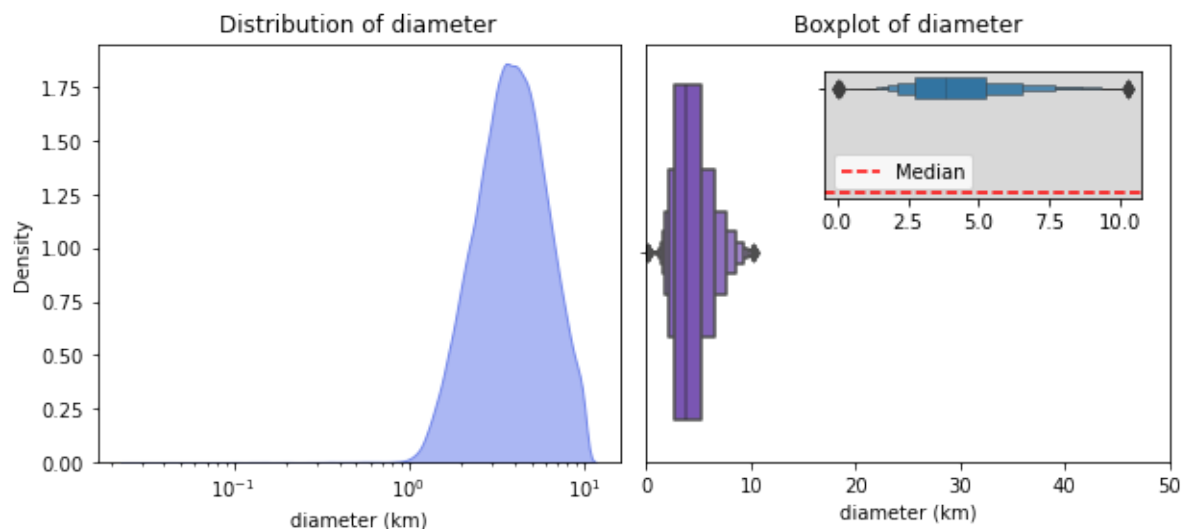
Out[94]: <AxesSubplot:>



```

In [98]: 1 fig = plt.figure(figsize=(10, 4))
2         plt.subplots_adjust(wspace=0.05)
3
4         plt.subplot(121, xlabel='diameter (km)', title='Distribution of diameter')
5         sns.kdeplot(data=df, x='diameter', fill=True, log_scale=True, color='#7387
6
7         plt.subplot(122, title='Boxplot of diameter', xlim=(0, 50))
8         sns.boxenplot(data=df, x='diameter', color='#7647C2')
9         plt.gca().set(xlabel='diameter (km)')
10
11        fig.add_axes([0.65, 0.6, 0.23, 0.23])
12        zoom_out_ax = sns.boxenplot(data=df, x='diameter', linewidth=.5)
13        zoom_out_ax.set_facecolor('#D8D8D8')
14        plt.xlabel('')
15
16        # Add median line to the boxplot
17        median_value = df['diameter'].median()
18        plt.axhline(y=median_value, color='red', linestyle='--', label='Median')
19        plt.legend()
20
21        plt.show()

```



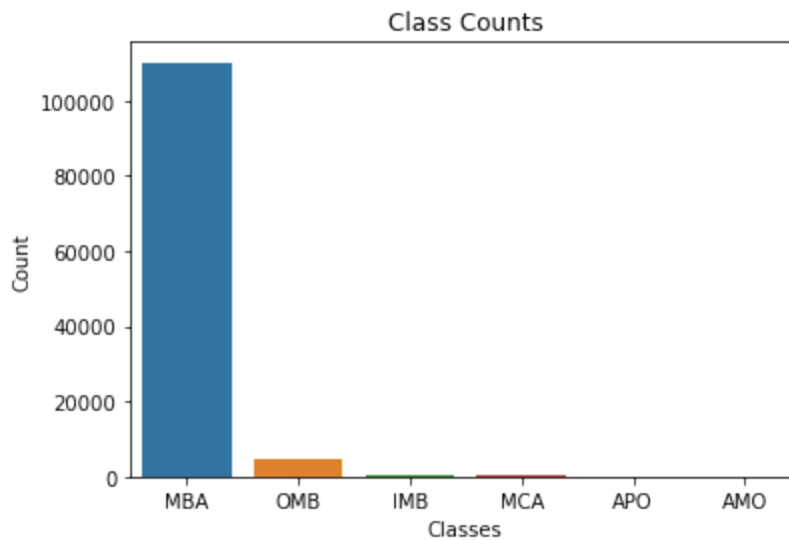
Alright, we see that the dataset contains values from 0.0025 to 10.0+, with the most of the values being in a range from 1 to 10. Even though some asteroids seem to be huge and have really large diameters of 10 km or more, the median diameter for the asteroids from this dataset is around 3.8 km.

Note, it's wrong to think that the greater the diameter of the asteroid, the more dangerous it is for Earth. The asteroid is considered potentially hazardous only if it is a Near-Earth object, and generally largest asteroids are less common among NEAs, because they have a greater gravitational attraction to the Sun, which causes them to be more stable in their orbits farther from Earth. Additionally, larger asteroids are more rare overall, as they represent a smaller fraction of the total population of asteroids in our solar system.


```
In [115]: 1 df['classes'].value_counts()
```

```
Out[115]: MBA      110087  
OMB       4930  
IMB       372  
MCA       244  
APO       112  
AMO        80  
Name: classes, dtype: int64
```

```
In [114]: 1 class_counts = df['classes'].value_counts().to_dict()  
2 class_counts = df['classes'].value_counts()  
3 sns.barplot(x=class_counts.index, y=class_counts.values)  
4 plt.xlabel('Classes')  
5 plt.ylabel('Count')  
6 plt.title('Class Counts')  
7 plt.show()
```



We see that MBA is certainly leading with around 92% of all asteroids being of that class. Our dataset contains some asteroids from the Outer Main-belt, and much less of other types. Not to mention, that we have only 80 asteroids orbiting outside AMO, all other asteroids usually cross the plane of orbit of certain planets in the solar system.

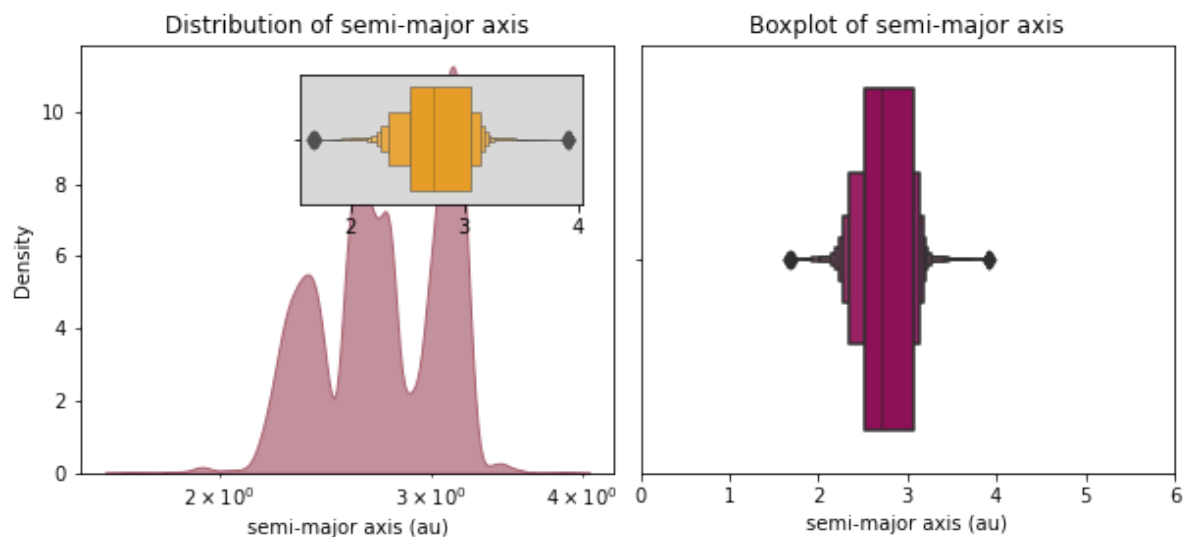
This actually might make it difficult for our model to generalize when it comes down to predicting diameters in the long run. We would want to make the model predict the diameters for far asteroids correctly, but as of now we won't focus on it, since most of the asteroids we have are MBAs.

```

In [116]: 1 fig = plt.figure(figsize=(10,4))
2 plt.subplots_adjust(wspace=0.05)
3
4 plt.subplot(121, xlabel='semi-major axis (au)', title='Distribution of sem
5 sns.kdeplot(data=df, x='a', fill=True, log_scale=True, color='#aa6373', al
6
7 plt.subplot(122, title='Boxplot of semi-major axis', xlim=(0,6))
8 sns.boxenplot(data=df, x='a', saturation=.8, color='#9e0059')
9 plt.gca().set(xlabel='semi-major axis (au)')
10
11 fig.add_axes([0.28,0.6,0.20,0.23])
12 zoom_out_ax = sns.boxenplot(data=df, x='a', saturation=.8, color='#fca311'
13 zoom_out_ax.set_facecolor('#D8D8D8')
14 plt.xlabel('')

```

Out[116]: Text(0.5, 0, '')



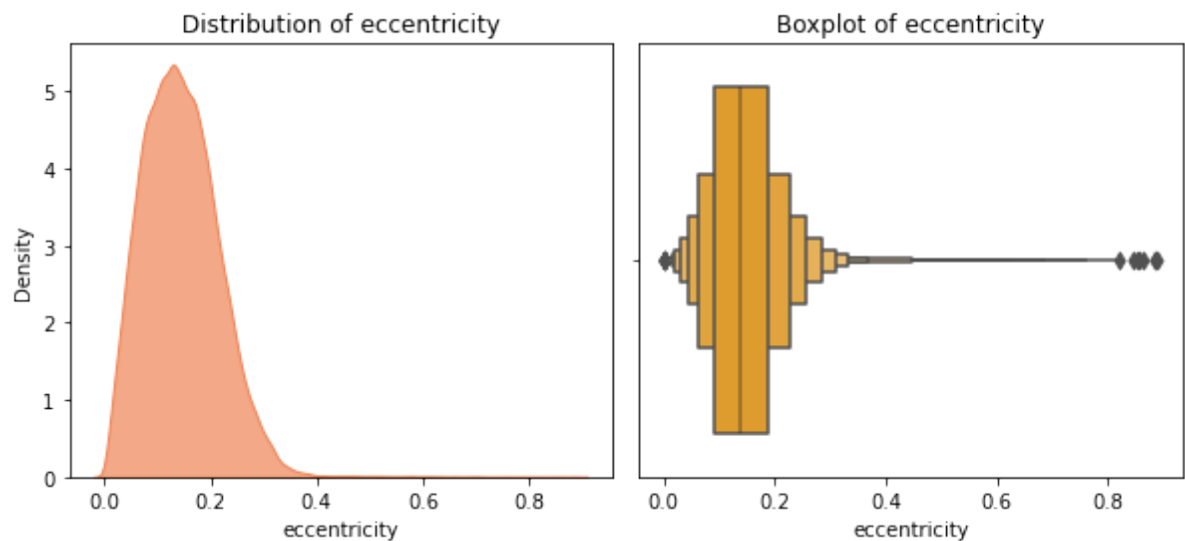
So we see that the most popular value for semi-major axis is around 2.5 and 3, and, overall, it ranges between 1.5 au and 4 au, depending on the class of the asteroid.

```

In [118]: 1 fig = plt.figure(figsize=(10,4))
          2 plt.subplots_adjust(wspace=0.05)
          3
          4 plt.subplot(121, xlabel='eccentricity', title='Distribution of eccentricit
          5 sns.kdeplot(data=df, x='e',fill=True, color='#ef8354', alpha=.7, )
          6
          7 plt.subplot(122, title='Boxplot of eccentricity')
          8 sns.boxenplot(data=df, x='e',color='#fca311')
          9 plt.gca().set(xlabel='eccentricity',)

```

Out[118]: [Text(0.5, 0, 'eccentricity')]



The eccentricity ranges from 0.0 to 0.3, which says that most asteroids have slightly elongated orbit.

```

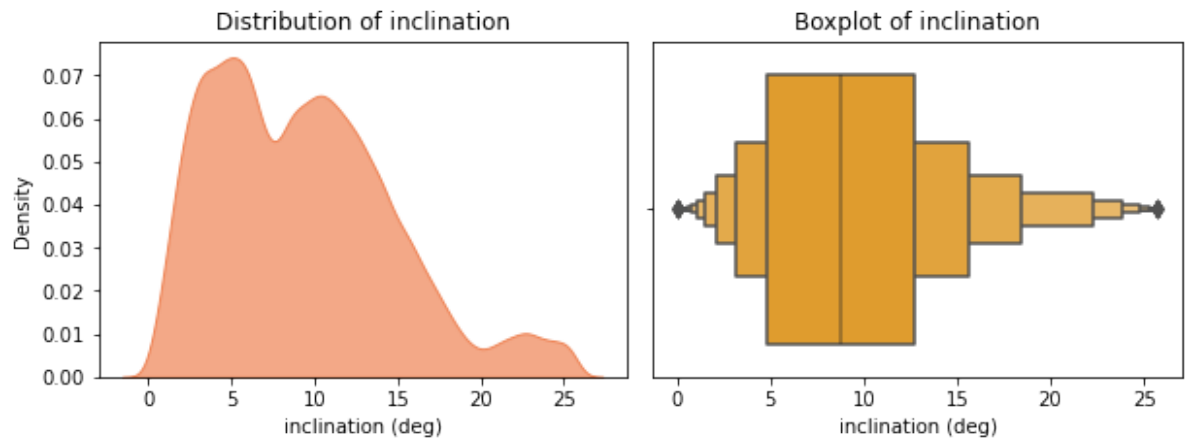
In [128]: 1 fig = plt.figure(figsize=(10,7))
          2 plt.subplots_adjust(wspace=0.05)
          3
          4 plt.subplot(221, xlabel='inclination (deg)', title='Distribution of inclin
          5 sns.kdeplot(data=df, x='i',fill=True, color='#ef8354', alpha=.7, )
          6
          7 plt.subplot(222)
          8 sns.boxenplot(data=df, x='i',color='#fca311')
          9 plt.gca().set(xlabel='inclination (deg)', title='Boxplot of inclination')

```

```

Out[128]: [Text(0.5, 0, 'inclination (deg)'), Text(0.5, 1.0, 'Boxplot of inclination')]

```



So we see that most of the eccentricities are from 0 deg to 15 deg. This may be a hint telling us that most of the asteroids in our dataset are MBAs (low eccentricity due to stable orbit). Not to mention that for these values of inclination we have the range for semi-major axis from 1.5 au to 4 au. Such semi-major axes commonly have MBAs, OMBs, IMBs, APOs, MCAs, APU.

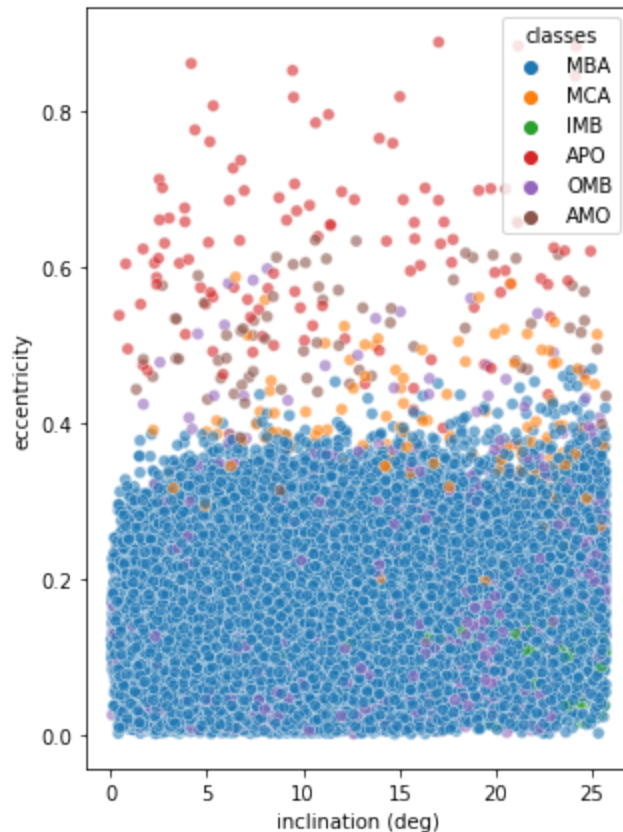
Therefore, the relationship between an asteroid's inclination and eccentricity can provide important information about the class asteroid belongs to, and thus, its overall characteristics

```
In [133]: 1 fig = plt.figure(figsize=(10,7))
2 plt.subplots_adjust(wspace=0.05)
3
4 plt.subplot(2,2,(1,3), xlabel='inclination (deg)', ylabel='eccentricity')
5 sns.scatterplot(data=df, x='i', y='e', hue='classes',
6                 alpha=.6)
```

Out[133]: <AxesSubplot:xlabel='inclination (deg)', ylabel='eccentricity'>

c:\Users\MEER\anaconda3\lib\site-packages\IPython\core\pylabtools.py:132: Use
rWarning: Creating legend with loc="best" can be slow with large amounts of data.

fig.canvas.print_figure(bytes_io, **kw)



From the plots this can be concluded:

- MBAs typically have low eccentricities and inclinations, which means their orbits are relatively stable.
- MCAs usually have low-inclination but high-eccentricity orbits, because they may have been perturbed by the gravity of Mars.
- APOs have orbits that bring them close to or cross the orbit of Earth. As a result, they tend to have higher eccentricities and inclinations than MBAs. Certain types of NEAs such as ATEs tend to have higher inclinations than other.

```
In [22]: 1 from sklearn.model_selection import train_test_split, cross_val_score, GridS
2 from sklearn.pipeline import Pipeline
3 from sklearn.impute import SimpleImputer
4 from sklearn.ensemble import RandomForestRegressor
5 from sklearn.preprocessing import StandardScaler
6 from sklearn.tree import DecisionTreeRegressor
7 from sklearn.linear_model import LinearRegression
8 from sklearn.model_selection import train_test_split
9 from sklearn.metrics import mean_squared_error, r2_score
```

```
In [23]: 1 X = df.drop("diameter", axis=1)
2 y = df["diameter"]
```

```
In [24]: 1 # Horizontal Splitting
2 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, r
```

```
In [25]: 1 X_train.head()
```

```
Out[25]:
```

	e	a	i	ma	n	tp	moid	data_arc	r
52656	0.301162	2.876067	7.046305	264.388383	0.202072	2.459074e+06	1.007630	8770.0	
68924	0.082819	2.696773	4.710427	243.657783	0.222555	2.459123e+06	1.469140	7505.0	
56587	0.252213	2.563935	4.153450	297.359674	0.240073	2.458861e+06	0.906491	7325.0	
56679	0.191185	3.087673	12.786612	224.367310	0.181659	2.459347e+06	1.526220	7181.0	
118795	0.149940	2.036710	5.769140	74.209816	0.339086	2.456058e+06	0.720762	1570.0	

In [26]:

```

1
2
3 # Create a List of models to evaluate
4 models = [
5     ("Random Forest", RandomForestRegressor(random_state=42)),
6     ("Decision Tree", DecisionTreeRegressor(random_state=42)),
7     ("Linear Regression", LinearRegression())
8 ]
9
10 # Initiate best_model and its performance metrics
11 best_model = None
12 best_rmse = float('inf')
13
14 # Iterate over models and evaluate their performance
15 for name, model in models:
16     pipeline = Pipeline([
17         ("scaler", StandardScaler()),
18         ("model", model)
19     ])
20
21     # Fit the pipeline on training data
22     pipeline.fit(X_train, y_train)
23
24     # Make predictions on test data
25     y_pred = pipeline.predict(X_test)
26
27     # Calculate RMSE
28     rmse = mean_squared_error(y_test, y_pred, squared=False)
29
30     # Print the performance metric
31     print("Model:", name)
32     print("Test RMSE:", rmse)
33     print()
34
35     # Check if the current model has the best RMSE
36     if rmse < best_rmse:
37         best_rmse = rmse
38         best_model = pipeline
39
40 # Retrieve the best model
41 print("Best Model:", best_model)

```

Model: Random Forest
 Test RMSE: 0.5557912596909688

Model: Decision Tree
 Test RMSE: 0.8111603731138083

Model: Linear Regression
 Test RMSE: 0.8548016521438069

Best Model: Pipeline(steps=[('scaler', StandardScaler()),
 ('model', RandomForestRegressor(random_state=42))])

Hyperparameter tuning for decision tree regressor

```
In [19]: 1 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, r
2
3 X_train, X_val, y_train, y_val = train_test_split(
4     X_train, y_train, test_size = 0.2, random_state = 42
5 )
```

```
In [23]: 1 from sklearn.metrics import r2_score, mean_squared_error
2
3 model = make_pipeline(
4     DecisionTreeRegressor(random_state=42)
5 )
6 # Fit model to training data
7 model.fit(X_train, y_train)
8
9 # Make predictions on training and validation data
10 y_train_pred = model.predict(X_train)
11 y_val_pred = model.predict(X_val)
12
13 # Calculate R-squared for training and validation data
14 r2_train = r2_score(y_train, y_train_pred)
15 r2_val = r2_score(y_val, y_val_pred)
16
17 # Calculate RMSE for training and validation data
18 rmse_train = mean_squared_error(y_train, y_train_pred, squared=False)
19 rmse_val = mean_squared_error(y_val, y_val_pred, squared=False)
20
21 print("Training R-squared:", round(r2_train, 2))
22 print("Validation R-squared:", round(r2_val, 2))
23 print("Training RMSE:", round(rmse_train, 2))
24 print("Validation RMSE:", round(rmse_val, 2))
25
26 tree_depth = model.named_steps["decisiontreeregressor"].get_depth()
27 print("Tree Depth:", tree_depth)
```

Training R-squared: 1.0
 Validation R-squared: 0.78
 Training RMSE: 0.0
 Validation RMSE: 0.82
 Tree Depth: 34

```
In [22]: 1 depth_hyperparams = range(1,50,2)
```



```

In [24]: 1 from sklearn.metrics import r2_score, mean_squared_error
2
3 # Create empty lists for training R-squared and RMSE scores
4 training_r2_scores = []
5 validation_r2_scores = []
6 training_rmse_scores = []
7 validation_rmse_scores = []
8
9 for d in depth_hyperparams:
10     # Create model with `max_depth` of `d`
11     test_model = make_pipeline(
12         DecisionTreeRegressor(max_depth=d, random_state=42)
13     )
14     # Fit model to training data
15     test_model.fit(X_train, y_train)
16     # Make predictions on training and validation data
17     y_train_pred = test_model.predict(X_train)
18     y_val_pred = test_model.predict(X_val)
19     # Calculate R-squared scores
20     r2_train = r2_score(y_train, y_train_pred)
21     r2_val = r2_score(y_val, y_val_pred)
22     # Calculate RMSE scores
23     rmse_train = mean_squared_error(y_train, y_train_pred, squared=False)
24     rmse_val = mean_squared_error(y_val, y_val_pred, squared=False)
25     # Append scores to respective lists
26     training_r2_scores.append(r2_train)
27     validation_r2_scores.append(r2_val)
28     training_rmse_scores.append(rmse_train)
29     validation_rmse_scores.append(rmse_val)
30
31 print("Training R-squared Scores:", training_r2_scores[:3])
32 print("Validation R-squared Scores:", validation_r2_scores[:3])
33 print("Training RMSE Scores:", training_rmse_scores[:3])
34 print("Validation RMSE Scores:", validation_rmse_scores[:3])

```


Training R-squared Scores: [0.20600611481076558, 0.4889495215446231, 0.7123824086864605]

Validation R-squared Scores: [0.2002173232274237, 0.48201726136407086, 0.7011513634239879]

Training RMSE Scores: [1.5637730185136285, 1.254576176135037, 0.9411803666404298]

Validation RMSE Scores: [1.5662154831007629, 1.2604428531930125, 0.957395001650114]

```
In [25]: 1 # Plot `depth_hyperparams`, `training_acc`  
2 plt.plot(depth_hyperparams, training_r2_scores, label = "Training")  
3 plt.plot(depth_hyperparams, validation_r2_scores, label = "Validation")  
4 plt.xlabel("Max_Depth")  
5 plt.ylabel("Accuracy_Score")  
6 plt.legend();
```



Model accuracy after Tuning

```

In [26]: 1 from sklearn.metrics import r2_score, mean_squared_error
2
3
4 model = make_pipeline(
5     DecisionTreeRegressor(max_depth=10, random_state=42)
6 )
7 # Fit model to training data
8 model.fit(X_train, y_train)
9
10 # Make predictions on training and validation data
11 y_train_pred = model.predict(X_train)
12 y_test_pred = model.predict(X_test)
13
14 # Calculate R-squared for training and validation data
15 r2_train = r2_score(y_train, y_train_pred)
16 r2_test = r2_score(y_test, y_test_pred)
17
18 # Calculate RMSE for training and validation data
19 rmse_train = mean_squared_error(y_train, y_train_pred, squared=False)
20 rmse_test = mean_squared_error(y_test, y_test_pred, squared=False)
21
22 print("Training R-squared:", round(r2_train, 2))
23 print("Validation R-squared:", round(r2_test, 2))
24 print("Training RMSE:", round(rmse_train, 2))
25 print("Validation RMSE:", round(rmse_test, 2))
26
27 tree_depth = model.named_steps["decisiontreeregressor"].get_depth()
28 print("Tree Depth:", tree_depth)

```

Training R-squared: 0.88
 Validation R-squared: 0.84
 Training RMSE: 0.61
 Validation RMSE: 0.71
 Tree Depth: 10

```

In [ ]: 1 features = X_train.columns
2 importances = model.named_steps["decisiontreeregressor"].feature_importances_
3
4 print("Features:", features[:3])
5 print("Importances:", importances[:3])

```

```

In [ ]: 1 feat_imp = pd.Series(importances, index=features).sort_values()
2 feat_imp.head()

```

```

In [ ]: 1 # Create horizontal bar chart
2 feat_imp.plot(kind = "barh")
3 plt.xlabel("Gini Importances")
4 plt.ylabel("Features");

```

Hyperparameter tuning for random forest regressor

```
In [27]: 1 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, r
2
3 X_train, X_val, y_train, y_val = train_test_split(
4     X_train, y_train, test_size = 0.2, random_state = 42
5 )
```

```
In [28]: 1 from sklearn.metrics import r2_score, mean_squared_error
2
3 model = make_pipeline(
4     RandomForestRegressor(random_state=42)
5 )
6 # Fit model to training data
7 model.fit(X_train, y_train)
8
9 # Make predictions on training and validation data
10 y_train_pred = model.predict(X_train)
11 y_val_pred = model.predict(X_val)
12
13 # Calculate R-squared for training and validation data
14 r2_train = r2_score(y_train, y_train_pred)
15 r2_val = r2_score(y_val, y_val_pred)
16
17 # Calculate RMSE for training and validation data
18 rmse_train = mean_squared_error(y_train, y_train_pred, squared=False)
19 rmse_val = mean_squared_error(y_val, y_val_pred, squared=False)
20
21 print("Training R-squared:", round(r2_train, 2))
22 print("Validation R-squared:", round(r2_val, 2))
23 print("Training RMSE:", round(rmse_train, 2))
24 print("Validation RMSE:", round(rmse_val, 2))
```

Training R-squared: 0.99

Validation R-squared: 0.9

Training RMSE: 0.21

Validation RMSE: 0.55

```
-----
AttributeError                                Traceback (most recent call last)
<ipython-input-28-9b3be2bbcbcb> in <module>
    24 print("Validation RMSE:", round(rmse_val, 2))
    25
----> 26 tree_depth = model.named_steps["randomforestregressor"].get_depth()
    27 print("Tree Depth:", tree_depth)
```

AttributeError: 'RandomForestRegressor' object has no attribute 'get_depth'

```
In [29]: 1 forest = model.named_steps["randomforestregressor"]
          2 tree_depths = [estimator.get_depth() for estimator in forest.estimators_]
          3 print("Tree Depths:", tree_depths)
```

```
Tree Depths: [34, 34, 34, 35, 34, 34, 34, 36, 34, 33, 34, 35, 36, 34, 36, 36,
35, 37, 35, 34, 32, 35, 35, 34, 32, 35, 35, 34, 32, 39, 33, 40, 37, 33, 31, 3
6, 33, 34, 39, 33, 32, 33, 38, 33, 38, 36, 37, 35, 34, 34, 39, 39, 35, 36, 3
4, 37, 34, 34, 32, 34, 34, 36, 40, 35, 34, 34, 38, 35, 33, 34, 33, 33, 35, 3
6, 32, 35, 34, 36, 35, 34, 35, 36, 38, 34, 36, 35, 34, 34, 35, 35, 35, 36, 3
8, 36, 33, 37, 34, 35, 35, 34]
```

```
In [30]: 1 depth_hyperparams = range(1,70,2)
```

```

In [31]: 1 from sklearn.metrics import r2_score, mean_squared_error
2
3 # Create empty lists for training R-squared and RMSE scores
4 training_r2_scores = []
5 validation_r2_scores = []
6 training_rmse_scores = []
7 validation_rmse_scores = []
8
9 for d in depth_hyperparams:
10     # Create model with `max_depth` of `d`
11     test_model = make_pipeline(
12         RandomForestRegressor(max_depth=d, random_state=42)
13     )
14     # Fit model to training data
15     test_model.fit(X_train, y_train)
16     # Make predictions on training and validation data
17     y_train_pred = test_model.predict(X_train)
18     y_val_pred = test_model.predict(X_val)
19     # Calculate R-squared scores
20     r2_train = r2_score(y_train, y_train_pred)
21     r2_val = r2_score(y_val, y_val_pred)
22     # Calculate RMSE scores
23     rmse_train = mean_squared_error(y_train, y_train_pred, squared=False)
24     rmse_val = mean_squared_error(y_val, y_val_pred, squared=False)
25     # Append scores to respective lists
26     training_r2_scores.append(r2_train)
27     validation_r2_scores.append(r2_val)
28     training_rmse_scores.append(rmse_train)
29     validation_rmse_scores.append(rmse_val)
30
31 print("Training R-squared Scores:", training_r2_scores[:3])
32 print("Validation R-squared Scores:", validation_r2_scores[:3])
33 print("Training RMSE Scores:", training_rmse_scores[:3])
34 print("Validation RMSE Scores:", validation_rmse_scores[:3])

```


Training R-squared Scores: [0.20933468097939467, 0.5281564120002881, 0.752907486313424]

Validation R-squared Scores: [0.20461740780090842, 0.5240100697172042, 0.7446712121419738]

Training RMSE Scores: [1.5604917661449513, 1.2054915359815777, 0.8723583660226655]

Validation RMSE Scores: [1.5619011952993584, 1.2082711306927323, 0.8849432238198869]

```
In [32]: 1 # Plot `depth_hyperparams`, `training_acc`  
2 plt.plot(depth_hyperparams, training_r2_scores, label = "Training")  
3 plt.plot(depth_hyperparams, validation_r2_scores, label = "Validation")  
4 plt.xlabel("Max_Depth")  
5 plt.ylabel("Accuracy_Score")  
6 plt.legend();
```



```
In [13]: 1 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, r
```

```
In [14]: 1 X_train.shape
```

```
Out[14]: (16000, 15)
```

```
In [41]: 1 from sklearn.metrics import r2_score, mean_squared_error
2
3 model = make_pipeline(
4     RandomForestRegressor(max_depth=12, random_state=42)
5 )
6 # Fit model to training data
7 model.fit(X_train, y_train)
8
9 # Make predictions on training and validation data
10 y_train_pred = model.predict(X_train)
11 y_test_pred = model.predict(X_test)
12
13 # Calculate R-squared for training and validation data
14 r2_train = r2_score(y_train, y_train_pred)
15 r2_test = r2_score(y_test, y_test_pred)
16
17 # Calculate RMSE for training and validation data
18 rmse_train = mean_squared_error(y_train, y_train_pred, squared=False)
19 rmse_test = mean_squared_error(y_test, y_test_pred, squared=False)
20
21 print("Training R-squared:", round(r2_train, 2))
22 print("Validation R-squared:", round(r2_test, 2))
23 print("Training RMSE:", round(rmse_train, 2))
24 print("Validation RMSE:", round(rmse_test, 2))
```

Training R-squared: 0.93
Validation R-squared: 0.89
Training RMSE: 0.45
Validation RMSE: 0.58

```
In [46]: 1 print("Training R-squared:", round(r2_train, 2))
2 print("Test R-squared:", round(r2_test, 2))
3 print("Training RMSE:", round(rmse_train, 2))
4 print("Test RMSE:", round(rmse_test, 2))
```


Training R-squared: 0.93
Test R-squared: 0.89
Training RMSE: 0.45
Test RMSE: 0.58

```
In [44]: 1 import pickle
2 pickle.dump(model, open('model.pkl', 'wb'))
3
```

```
In [15]: 1 import pickle
2 model = pickle.load(open('model.pkl', 'rb'))
```




```
In [16]: 1 import shap
2
3 # Step 2: Identify the top three features
4 importance_scores = model.named_steps["randomforestregressor"].feature_imp
5 top_features_indices = importance_scores.argsort()[-3:][::-1]
6 top_features = X_train.columns[top_features_indices]
7
8 # Step 3: Calculate Shapley values
9 explainer = shap.Explainer(model.named_steps["randomforestregressor"])
10 shap_values = explainer.shap_values(X_test)
11
12 # Step 4: Generate Shapley plots for the top three features
13 shap.summary_plot(shap_values, X_test, feature_names=X_test.columns, plot_
```



```

In [17]: 1 import shap
2
3 # Calculate Shapley values
4 explainer = shap.Explainer(model.named_steps["randomforestregressor"])
5 shap_values = explainer.shap_values(X_test)
6
7 # Get feature importances
8 importance_scores = model.named_steps["randomforestregressor"].feature_imp
9
10 # Sort features based on importance scores
11 top_features_indices = importance_scores.argsort()[::-1]
12 top_features = X_train.columns[top_features_indices]
13
14 # Generate Shapley plots for increasing diameter prediction (top three fea
15 shap.summary_plot(shap_values, X_test, feature_names=X_test.columns, plot_
16
17 # Find the top feature for decreasing diameter prediction
18 top_feature_decreasing = top_features[0]
19 print("Top feature for decreasing diameter prediction:", top_feature_decre

```



Top feature for decreasing diameter prediction: n_obs_used

In []:

1	
---	--